

## Import Dependencies

```
In [ ]: import pandas as pd
        from sklearn.model_selection import train_test_split
        import re
        import nltk
        from nltk.tokenize import word_tokenize
        from collections import Counter
        from keras.preprocessing.sequence import pad_sequences
        import matplotlib.pyplot as plt
        import numpy as np
        import torch
        import torch.nn as nn
        from torch.utils.data import DataLoader, TensorDataset
        import torch.optim as optim
        import torch.nn.functional as F
```

```
In [ ]: class SentimentClassifier(nn.Module):
        def __init__(self, vocab_size, embedding_dim, hidden_dim, output_dim):
            super(SentimentClassifier, self).__init__()
            self.embedding = nn.Embedding(vocab_size, embedding_dim)
            self.fc = nn.Linear(embedding_dim, hidden_dim)
            self.relu = nn.ReLU()
            self.out = nn.Linear(hidden_dim, output_dim)
            self.sigmoid = nn.Sigmoid()

        def forward(self, x):
            embedded = self.embedding(x)
            hidden = self.fc(embedded.mean(dim=1))
            hidden = self.relu(hidden)
            output = self.out(hidden)
            return self.sigmoid(output)
```

```
In [ ]: # Load dataset
        df = pd.read_csv('./imdb.csv')

        # Splitting the dataset into training, validation, and testing sets
        train_df, temp_df = train_test_split(df, test_size=0.3, random_state=42) #
        val_df, test_df = train_test_split(temp_df, test_size=0.5, random_state=42)

        # Cleaning Data
        def clean_text(text):
            # Add whitespace around punctuation
            text = re.sub(r'([.,!()?])', r' \1 ', text)
            # Remove non-punctuation symbols
            text = re.sub(r'^a-zA-Z.,!()? ]', '', text)
            return text.strip()

        # Apply cleaning to all splits
        train_df['cleaned_text'] = train_df['review'].apply(clean_text)
        val_df['cleaned_text'] = val_df['review'].apply(clean_text)
        test_df['cleaned_text'] = test_df['review'].apply(clean_text)
```

```
In [ ]: # Map string labels to numeric values: 'negative' to 0 and 'positive' to 1
label_map = {'negative': 0, 'positive': 1}
train_df['sentiment'] = train_df['sentiment'].map(label_map).astype(float)
val_df['sentiment'] = val_df['sentiment'].map(label_map).astype(float)
test_df['sentiment'] = test_df['sentiment'].map(label_map).astype(float)

# Download NLTK tokenizer model
nltk.download('punkt')

# Tokenization
tokenized_reviews = [word_tokenize(review.lower()) for review in train_df['cleaned_review']]
val_tokenized = [word_tokenize(review.lower()) for review in val_df['cleaned_review']]
test_tokenized = [word_tokenize(review.lower()) for review in test_df['cleaned_review']]

# Build Vocabulary
word_counts = Counter(word for review in tokenized_reviews for word in review)
vocabulary = {word: i + 1 for i, (word, _) in enumerate(word_counts.most_common())}

# Convert Text to Integer Sequences
reviews_int = [[vocabulary[word] for word in review if word in vocabulary] for review in tokenized_reviews]
val_reviews_int = [[vocabulary[word] for word in review if word in vocabulary] for review in val_tokenized]
test_reviews_int = [[vocabulary[word] for word in review if word in vocabulary] for review in test_tokenized]
review_lengths = [len(review) for review in tokenized_reviews]

# Find a suitable max length
percentile = 95
max_len = int(np.percentile(review_lengths, percentile))

# Padding Sequences
padded_reviews = pad_sequences(reviews_int, maxlen=max_len, padding='post', dtype='int')
val_padded_reviews = pad_sequences(val_reviews_int, maxlen=max_len, padding='post', dtype='int')
test_padded_reviews = pad_sequences(test_reviews_int, maxlen=max_len, padding='post', dtype='int')
```

[nltk\_data] Downloading package punkt to /Users/nickparov/nltk\_data...  
[nltk\_data] Package punkt is already up-to-date!

```
In [ ]: # Parameters
vocab_size = 10000 + 1 # +1 for padding token
embedding_dim = 100 # common choice for embedding dimension
hidden_dim = 64 # mid-range for hidden dimension
output_dim = 1 # binary

# init model
model = SentimentClassifier(vocab_size, embedding_dim, hidden_dim, output_dim)

# Parameters
batch_size = 32

train_inputs = torch.tensor(padded_reviews, dtype=torch.long)
val_inputs = torch.tensor(val_padded_reviews, dtype=torch.long)
test_inputs = torch.tensor(test_padded_reviews, dtype=torch.long)

# labels
train_labels = torch.tensor(train_df['sentiment'].values, dtype=torch.float32)
val_labels = torch.tensor(val_df['sentiment'].values, dtype=torch.float32)
```

```

test_labels = torch.tensor(test_df['sentiment'].values, dtype=torch.float32)

# Train DataLoader
train_data = TensorDataset(train_inputs, train_labels)
train_loader = DataLoader(train_data, batch_size=batch_size, shuffle=True)

# Validation DataLoader
val_data = TensorDataset(val_inputs, val_labels)
val_loader = DataLoader(val_data, batch_size=batch_size, shuffle=True)

# Test DataLoader
test_data = TensorDataset(test_inputs, test_labels)
test_loader = DataLoader(test_data, batch_size=batch_size, shuffle=False)

```

```

In [ ]: # Define optimizer and loss function
optimizer = optim.Adam(model.parameters(), lr=0.001)
criterion = nn.BCELoss()

num_epochs = 8

for epoch in range(num_epochs):
    # Training Phase
    model.train()
    total_loss, total, correct = 0, 0, 0

    for inputs, labels in train_loader:
        optimizer.zero_grad()
        outputs = model(inputs)
        loss = criterion(outputs.squeeze(), labels)
        loss.backward()
        optimizer.step()

        total_loss += loss.item()
        predicted = (outputs.squeeze() > 0.5).float()
        correct += (predicted == labels).sum().item()
        total += labels.size(0)

    train_loss = total_loss / len(train_loader)
    train_accuracy = correct / total
    print(f'Epoch {epoch+1}/{num_epochs}, Train Loss: {train_loss:.4f}, Trai

    # Validation Phase
    model.eval()
    val_loss, val_correct, val_total = 0, 0, 0
    with torch.no_grad():
        for inputs, labels in val_loader:
            outputs = model(inputs)
            loss = criterion(outputs.squeeze(), labels)
            val_loss += loss.item()
            predicted = (outputs.squeeze() > 0.5).float()
            val_correct += (predicted == labels).sum().item()
            val_total += labels.size(0)

    val_loss /= len(val_loader)

```

```
val_accuracy = val_correct / val_total
print(f'Epoch {epoch+1}/{num_epochs}, Validation Loss: {val_loss:.4f}, V
```

```
Epoch 1/8, Train Loss: 0.5962, Train Accuracy: 0.6543
Epoch 1/8, Validation Loss: 0.3987, Validation Accuracy: 0.8391
Epoch 2/8, Train Loss: 0.3382, Train Accuracy: 0.8609
Epoch 2/8, Validation Loss: 0.3087, Validation Accuracy: 0.8735
Epoch 3/8, Train Loss: 0.2734, Train Accuracy: 0.8897
Epoch 3/8, Validation Loss: 0.2917, Validation Accuracy: 0.8751
Epoch 4/8, Train Loss: 0.2434, Train Accuracy: 0.9019
Epoch 4/8, Validation Loss: 0.2644, Validation Accuracy: 0.8925
Epoch 5/8, Train Loss: 0.2209, Train Accuracy: 0.9130
Epoch 5/8, Validation Loss: 0.2735, Validation Accuracy: 0.8865
Epoch 6/8, Train Loss: 0.2054, Train Accuracy: 0.9206
Epoch 6/8, Validation Loss: 0.3199, Validation Accuracy: 0.8607
Epoch 7/8, Train Loss: 0.1924, Train Accuracy: 0.9247
Epoch 7/8, Validation Loss: 0.2559, Validation Accuracy: 0.8967
Epoch 8/8, Train Loss: 0.1808, Train Accuracy: 0.9301
Epoch 8/8, Validation Loss: 0.2655, Validation Accuracy: 0.8944
```

```
In [ ]: # Model Testing
model.eval()
test_loss, test_correct, test_total = 0, 0, 0

with torch.no_grad():
    for inputs, labels in test_loader:
        outputs = model(inputs)
        loss = criterion(outputs.squeeze(), labels)
        test_loss += loss.item()
        predicted = (outputs.squeeze() > 0.5).float()
        test_correct += (predicted == labels).sum().item()
        test_total += labels.size(0)

test_loss /= len(test_loader)
test_accuracy = test_correct / test_total
print(f'Test Loss: {test_loss:.4f}, Test Accuracy: {test_accuracy:.4f}')

# Extract the embedding layer weights
embedding_weights = model.embedding.weight.data

# Calculate the influence of each word
word_influence = torch.norm(embedding_weights, dim=1)

# Create a dictionary of words and their corresponding influence
word_influence_dict = {word: influence.item() for word, influence in zip(voc

# Sort words by influence
sorted_words = sorted(word_influence_dict.items(), key=lambda x: x[1], rever

# Get top 20 influential words for positive and negative reviews
# higher values indicate positive influence and lower values indicate negati
top_positive_words = sorted_words[:20]
top_negative_words = sorted_words[-20:]

# Print top 20 influential words for positive reviews
```

```

print("Top 20 Positive Words:")
for word, influence in top_positive_words[:20]:
    print(f"{word}: {influence}")

# Print top 20 influential words for negative reviews
print("\nTop 20 Negative Words:")
for word, influence in top_negative_words[:20]:
    print(f"{word}: {influence}")

```

Test Loss: 0.2616, Test Accuracy: 0.9005

Top 20 Positive Words:

- entire: 26.370708465576172
- kind: 25.498929977416992
- camera: 23.64904022216797
- credits: 21.489665985107422
- men: 20.365657806396484
- appropriate: 20.28223419189453
- page: 20.166799545288086
- giving: 20.06629753112793
- views: 19.575176239013672
- green: 19.46065902709961
- boyfriend: 19.23826026916504
- mysterious: 19.13362693786621
- beautiful: 19.074342727661133
- weapons: 18.98100471496582
- wasted: 18.938899993896484
- mean: 18.158227920532227
- sounds: 18.05463218688965
- loves: 17.826236724853516
- sea: 17.820823669433594
- insane: 17.72930908203125

Top 20 Negative Words:

- sunday: 8.13501262664795
- producers: 8.122499465942383
- marshall: 8.121447563171387
- made: 8.105531692504883
- washing: 8.09601879119873
- stupid: 8.087364196777344
- accidentally: 8.057730674743652
- enemy: 8.051525115966797
- sigh: 8.036587715148926
- newcomer: 8.004074096679688
- shows: 7.994915008544922
- june: 7.987314224243164
- symphony: 7.984182357788086
- economic: 7.974999904632568
- thrill: 7.962036609649658
- alexandra: 7.954861640930176
- misses: 7.931981086730957
- worries: 7.739133834838867
- multiple: 7.527898788452148
- predator: 6.8740081787109375

In [ ]: